

AWS Developer Playbook — Expanded

Audience: Developers, architects and DevOps teams building cloud-native applications for an asset-management / financial-services client

Scope: Practical developer playbook per service (S3 notifications, SQS, SNS, Lambda, DynamoDB, CloudWatch Logs, CloudWatch Alarms, Step Functions, EventBridge, EC2, ECS). Each service includes: overview, architecture patterns, best practices, example IAM policy, CDK (TypeScript) snippet, sample retry/deduplication code, observability, security / regulatory considerations, DR strategies, testing & CI/CD, and a short runbook for common operational failures.

How to use this doc: Copy relevant sections into your internal playbooks, adapt IAM policies & CDK to your account IDs and resource names, and store secrets/keys in Secrets Manager/SSM. This document assumes AWS CDK v2 / TypeScript examples. If you prefer Python CDK or full CloudFormation, tell me and I'll generate it.

Table of contents

1. S3 Event Notification
2. SQS
3. SNS
4. Lambda
5. DynamoDB
6. CloudWatch Logs
7. CloudWatch Alarms
8. Step Functions
9. EventBridge
10. EC2
11. ECS

1. S3 Event Notification

Overview

S3 event notifications can publish `ObjectCreated`, `ObjectRemoved`, and other events to SQS, SNS, or Lambda. Delivery is **at-least-once** and occasionally events can be delayed, duplicated, or (rarely) lost.

Developer playbook & patterns

- Treat S3 events as advisory: the object store is the source-of-truth. Implement periodic reconciliation (list/compare) for completeness.
- Use `versionId` for deduping and to detect updates vs deletes.
- For large objects / multipart uploads use `s3:ObjectCreated:CompleteMultipartUpload` to ensure full object availability.
- For heavy fan-out, prefer S3 -> EventBridge or S3 -> SNS -> SQS fan-out pattern to get durable buffering.
- For sensitive data, use S3 bucket policies + VPC endpoints to restrict access and avoid public egress.

Example IAM (least privilege for S3 to publish to SNS/SQS)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowS3ToSendMessage",
      "Effect": "Allow",
      "Principal": {"Service": "s3.amazonaws.com"},
      "Action": "sts:AssumeRole",
      "Resource": "*"
    }
  ]
}
```

Note: S3 uses a bucket notification configuration; most permission control is done via bucket policy granting `sns:Publish` or `sqs:SendMessage` to `arn:aws:iam::aws:policy/service-role/AmazonS3RoleForReplication` style principals as needed.

CDK snippet (TypeScript): S3 -> SQS notification

```
import * as cdk from 'aws-cdk-lib';
import { Stack, StackProps } from 'aws-cdk-lib';
import * as s3 from 'aws-cdk-lib/aws-s3';
import * as sqs from 'aws-cdk-lib/aws-sqs';
import * as s3n from 'aws-cdk-lib/aws-s3-notifications';

export class S3SqsStack extends Stack {
  constructor(scope: cdk.App, id: string, props?: StackProps) {
    super(scope, id, props);

    const bucket = new s3.Bucket(this, 'DataBucket', { versioned: true });
    const queue = new sqs.Queue(this, 'S3EventsQueue', { retentionPeriod:
cdk.Duration.days(14) });
```

```

        bucket.addEventNotification(s3.EventType.OBJECT_CREATED, new
s3n.SqsDestination(queue));
    }
}

```

Sample dedupe/idempotency code (Node.js lambda consumer)

```

// Use DynamoDB for dedupe store - table: s3-event-dedupe (pk: eventId)
const AWS = require('aws-sdk');
const ddb = new AWS.DynamoDB.DocumentClient();
const DEDUPE_TABLE = process.env.DEDUPE_TABLE;

exports.handler = async (event) => {
  for (const rec of event.Records) {
    const eventId = rec.eventID; // S3 eventID
    // Try conditional put: if not exists, process
    try {
      await ddb.put({
        TableName: DEDUPE_TABLE,
        Item: { eventId, processedAt: Date.now() },
        ConditionExpression: 'attribute_not_exists(eventId)'
      }).promise();
    } catch (e) {
      if (e.code === 'ConditionalCheckFailedException') {
        console.log('Duplicate event, skipping', eventId);
        continue;
      }
      throw e;
    }

    // process the object
    // ... business logic ...
  }
};

```

Observability

- Enable S3 server access logs and CloudTrail S3 data events for auditing.
- Emit structured logs in downstream processors referencing S3 `objectKey`, `versionId` and `eventId`.
- Track counts of events, duplicates, failed processings, and reconciliation differences.

Security & regulatory considerations

- Encryption: enforce `SSE-S3` or `SSE-KMS` for buckets with financial data. If KMS, use a CMK with tight key policy and key rotation.
- Bucket policies: restrict `PutObject` and `GetObject` to specific VPC endpoints or IAM principals. Require TLS (`aws:SecureTransport`).
- Data residency: create buckets in approved regions. Use Bucket Lock/Retention for immutable audit artifacts.
- Audit: enable CloudTrail data events for S3 to meet traceability requirements.

DR & Recovery

- Enable Cross-Region Replication (CRR) for critical data; monitor replication lag.
- Keep lifecycle snapshot backups (versioning + lifecycle policies + replication) for RPO.
- Reconciliation job: scheduled process that lists prefixes and compares to processed state (S3 list vs processed table) to find missing events and reprocess.

Runbook (common failures)

- Missing events: check S3 bucket notification configuration, CloudTrail data events and server access logs. Run reconciliation.
 - Duplicate processing: verify dedupe table TTL, conditional writes, exceptions during processing that may not have recorded processed marker.
-

2. SQS

Overview

Durable queuing for decoupling components. Standard queues: high throughput, at-least-once, unordered. FIFO queues: ordered and exactly-once semantics (within dedupe window) with lower throughput.

Developer playbook & patterns

- Use FIFO when order matters (trade flows, accounting ledgers) and throughput is compatible.
- Always implement idempotency on the consumer: use message dedupe-id or external dedupe store.
- For long processing: set `VisibilityTimeout` longer than expected processing time and extend with API if needed.
- Use DLQs to capture poison messages and fix/inspect them manually before reprocessing.
- For large payloads: store payload in S3 and put pointer in SQS (pointer pattern).

Example IAM (consumer role)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "sqs:ReceiveMessage",
        "sqs:DeleteMessage",
        "sqs:GetQueueAttributes",
        "sqs:ChangeMessageVisibility"
      ],
      "Resource": "arn:aws:sqs:REGION:ACCOUNT_ID:queue-name"
    }
  ]
}
```

CDK snippet: creating a FIFO queue with DLQ

```
import * as sqs from 'aws-cdk-lib/aws-sqs';

const dlq = new sqs.Queue(this, 'DLQ', { fifo: true, queueName: 'my-dlq.fifo' });
const queue = new sqs.Queue(this, 'MainQueue', {
  fifo: true,
  contentBasedDeduplication: false,
  queueName: 'my-main-queue.fifo',
  deadLetterQueue: { maxReceiveCount: 5, queue: dlq }
});
```

Sample consumer with retry/backoff and dedupe (Node.js + DynamoDB dedupe)

```
const { SQSClient, ReceiveMessageCommand, DeleteMessageCommand,
ChangeMessageVisibilityCommand } = require('@aws-sdk/client-sqs');
const ddb = new (require('@aws-sdk/client-dynamodb')).DynamoDBClient({});

// Example exponential backoff with jitter
function sleep(ms){return new Promise(r=>setTimeout(r,ms));}
function backoff(attempt){
  const base = 100; // ms
  const max = 10000; // ms
```

```

const exp = Math.min(max, base * Math.pow(2, attempt));
return Math.floor(exp / 2 + Math.random() * exp / 2); // decorrelated jitter
}

async function processMessage(msg){
  // dedupe via messageId or body hash
}

// consumer loop (pseudo)

```

Observability

- Metrics: ApproximateNumberOfMessagesVisible, ApproximateNumberOfMessagesNotVisible, ApproximateAgeOfOldestMessage, NumberOfMessagesSent/Received/Deleted, ApproximateNumberOfMessagesDelayed.
- Alert if oldest message age > SLA or visible message count increases beyond threshold.

Security & regulatory considerations

- Encryption: SSE-SQS (KMS) for message payload encryption; manage CMKs and key policies.
- Access control: restrict who can send to and receive from queues (cross-account topics restricted).
- Message retention: set retention period according to regulatory retention policies; use DLQ for forensics.

DR & Recovery

- Use cross-region replication patterns: publish events to a central bus or replicate critical data via S3/DynamoDB Global Tables; SQS itself is regional; for cross-region, use SNS/EventBridge or custom replication.
- Periodic snapshot of message pointers not straightforward; design idempotent replays using source-of-truth.

Runbook

- Increasing queue backlog: inspect consumer scaling (Lambda concurrency or EC2/ECS consumers), check for processing errors, check DLQ counts.
- Poison message loop: inspect DLQ, patch code to handle edge cases, then re-inject items safely.

3. SNS

Overview

Pub/sub service used for fan-out to multiple subscribers (HTTP/S, email, SMS, SQS, Lambda).

Developer playbook & patterns

- For guaranteed durable delivery to multiple subscribers, use SNS -> SQS for each subscriber (durable buffer) rather than direct HTTP.
- Use message attributes and subscription filtering to reduce recipient load.
- Protect topics with topic policies and require HTTPS endpoints or signed subscriptions for cross-account.

Example IAM policy (publish privilege)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "sns:Publish",
      "Resource": "arn:aws:sns:REGION:ACCOUNT_ID:topic-name"
    }
  ]
}
```

CDK snippet: SNS topic with SQS subscriptions

```
import * as sns from 'aws-cdk-lib/aws-sns';
import * as subs from 'aws-cdk-lib/aws-sns-subscriptions';

const topic = new sns.Topic(this, 'Topic');
const queue = new sqs.Queue(this, 'Queue');
topic.addSubscription(new subs.SqsSubscription(queue));
```

Observability

- Monitor `NumberOfMessagesPublished`, `NumberOfNotificationsDelivered`, `NumberOfNotificationsFailed`.
- Track retry counts and use CloudWatch metrics to trigger alerts for failures.

Security & regulatory

- Encryption: use KMS CMKs for encrypting messages.
- Contracts: message schemas should be documented and versioned.
- For SMS/email, respect regulatory consent and opt-out policies.

DR & Recovery

- Use durable subscribers (SQS) with retry and DLQ to enable replay.

4. Lambda

Overview

Serverless compute with maximum 15 minute runtime. Supports sync and async invocation models and event sources including S3, SQS, SNS, EventBridge.

Developer playbook & patterns

- Keep Lambdas small and single-purpose; prefer Step Functions for complex orchestration.
- Use versions and aliases for safe deployment; use traffic-shifting for canary/linear deployments.
- For low latency-critical functions use Provisioned Concurrency.
- Use Lambda Layers to share dependencies and reduce cold start size.

Example IAM execution role (minimal)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {"Effect": "Allow", "Action":
    ["logs:CreateLogGroup", "logs:CreateLogStream", "logs:PutLogEvents"], "Resource": "arn:aws:logs:REGION:ACCOUNT_ID:log-group:/aws/lambda/your-func*"},
    {"Effect": "Allow", "Action":
    ["dynamodb:PutItem", "dynamodb:GetItem"], "Resource": "arn:aws:dynamodb:REGION:ACCOUNT_ID:table/your-table"}
  ]
}
```

CDK snippet: Lambda with SQS event source

```
import * as lambda from 'aws-cdk-lib/aws-lambda';
import * as eventSources from 'aws-cdk-lib/aws-lambda-event-sources';

const fn = new lambda.Function(this, 'Processor', {
  runtime: lambda.Runtime.NODEJS_18_X,
  handler: 'index.handler',
  code: lambda.Code.fromAsset('lambda/processor'),
  timeout: cdk.Duration.seconds(60)
```



```
});
```

```
fn.addEventSource(new eventSources.SqsEventSource(queue, { batchSize: 10,  
maxBatchingWindow: cdk.Duration.seconds(30) }));
```

Sample retry/backoff pattern (Node.js) and async handling

- For **sync** invocations: return errors deliberately (HTTP 5xx) and rely on caller retry.
- For **async** invocations: use DLQ or Failure Destinations to capture failed events.

Node.js partial sample: extend SQS visibility timeout while processing

```
const AWS = require('aws-sdk');  
const sqs = new AWS.SQS();  
  
async function extendVisibility(queueUrl, receiptHandle, seconds) {  
    await sqs.changeMessageVisibility({ QueueUrl: queueUrl, ReceiptHandle:  
receiptHandle, VisibilityTimeout: seconds }).promise();  
}  
  
exports.handler = async (event) => {  
    for (const r of event.Records) {  
        const handle = r.receiptHandle; // attempt to extend as you near time limit  
        // periodically call extendVisibility(queueUrl, handle, 60);  
    }  
}
```

Observability

- Use structured logs and X-Ray for distributed tracing. Correlate logs by requestId and trace id.
- Monitor `Invocations`, `Errors`, `Duration`, `Throttles`, and `ConcurrentExecutions`.

Security & regulatory

- Avoid embedding secrets in environment variables unless encrypted; use Secrets Manager & IAM to access secrets at runtime.
- VPC placement: Lambda in VPC requires ENIs; avoid unless necessary. If in VPC, ensure access to necessary endpoints via VPC endpoints.
- For financial workloads, enable CloudTrail data events for Lambda to capture invocation/audit data.

DR & Recovery

- Use multiple regions for critical functions; deploy Canary aliases and traffic-shifting to test failover.
- Use DLQs for async failures and event archive (EventBridge) to replay events.

Runbook

- High error rate: check code errors in CloudWatch logs, check dependent services' availability, increase concurrency or add retries with jitter.

5. DynamoDB

Overview

NoSQL key-value and document store with single-digit ms latency. Design must be access-pattern driven.

Developer playbook & patterns

- Model for queries: collect all access patterns first, then design PK/SK and GSIs.
- Single-table design can reduce complexity and latency — document patterns and keep items small.
- Protect against hot partitions by salting keys or redesigning to distribute workload.

Example IAM policy for app to access table

```
{
  "Version": "2012-10-17",
  "Statement": [{"Effect": "Allow", "Action":
["dynamodb:GetItem", "dynamodb:PutItem", "dynamodb:Query", "dynamodb:UpdateItem"], "Resource": "arn:aws:dynamodb:us-east-1:123456789012:table/transactions"}]}
}
```

CDK snippet: table + stream + global secondary index

```
import * as dynamodb from 'aws-cdk-lib/aws-dynamodb';

const table = new dynamodb.Table(this, 'TxTable', {
  partitionKey: { name: 'pk', type: dynamodb.AttributeType.STRING },
  sortKey: { name: 'sk', type: dynamodb.AttributeType.STRING },
  billingMode: dynamodb.BillingMode.PAY_PER_REQUEST,
  stream: dynamodb.StreamViewType.NEW_AND_OLD_IMAGES
});

table.addGlobalSecondaryIndex({ indexName: 'GSI1', partitionKey: { name:
'gsi1pk', type: dynamodb.AttributeType.STRING }, projectionType:
dynamodb.ProjectionType.ALL });
```

Sample conditional write / idempotency (Node.js, AWS SDK v3)

```
import { DynamoDBClient, PutItemCommand } from '@aws-sdk/client-dynamodb';
const client = new DynamoDBClient({});

async function putIfNotExists(tableName, pk, sk, payload) {
  const params = {
    TableName: tableName,
    Item: { pk: { S: pk }, sk: { S: sk }, payload: { S:
JSON.stringify(payload) } },
    ConditionExpression: 'attribute_not_exists(pk)'
  };
  try { await client.send(new PutItemCommand(params)); return true; }
  catch (e) { if (e.name === 'ConditionalCheckFailedException') return false;
throw e; }
}
```

Observability

- Monitor throttled requests, read/write latency, consumed capacity. Set alarms on ThrottledRequests and high latency.
- Use CloudWatch Contributor Insights to spot hot keys.

Security & regulatory

- Use encryption at rest (AWS-managed or CMK). If CMK, control key usage via IAM and key policies.
- Use VPC endpoints for DynamoDB to avoid egress.
- For audit trails, enable CloudTrail for table-level API calls and Streams for CDC.

DR & Recovery

- Global Tables for multi-region active-active replication for low RTO/RPO.
- Enable Point-in-Time Recovery (PITR) for recovery to any second in 35 days.
- On-demand backups for retention beyond PITR windows.

Runbook

- Throttling: increase capacity (if provisioned) or enable autoscaling / move heavy reads to eventually-consistent reads or cache (DAX/Redis).
- Lost writes: check conditional write failures and transaction logs.

6. CloudWatch Logs

Overview

Central log store. Supports subscription filters, insights queries, and metric extraction.

Developer playbook & patterns

- Ship structured JSON logs with correlation IDs.
- Use log group naming convention: `/aws/<service>/<app>/<env>`.
- Limit verbose logs in production (set log level via environment variables and remote configuration).

Example IAM to push logs

```
{
  "Effect": "Allow",
  "Action": ["logs:CreateLogGroup", "logs:CreateLogStream", "logs:PutLogEvents"],
  "Resource": "arn:aws:logs:REGION:ACCOUNT_ID:log-group:/aws/lambda/*"
}
```

CDK snippet: log group + retention

```
import * as logs from 'aws-cdk-lib/aws-logs';

new logs.LogGroup(this, 'AppLogGroup', { logGroupName: '/aws/app/prod',
retention: logs.RetentionDays.ONE_YEAR, removalPolicy:
cdk.RemovalPolicy.RETAIN });
```

Observability & metric filters

- Create metric filters for critical events (authentication failures, exceptions) and route to CloudWatch Alarms.
- Use CloudWatch Logs Insights for ad-hoc investigations.

Security & regulatory

- Enable encryption for log groups if logs contain sensitive data.
- Ensure log retention policies comply with regulatory retention requirements; export to secure archival (S3 with immutability) if required.

DR & Recovery

- Export critical logs to S3 (lifecycle + Glacier) for long-term retention and legal hold.
-

7. CloudWatch Alarms

Overview

Alerting on metric thresholds and composite conditions.

Developer playbook & patterns

- Use multiple datapoints and evaluation periods to avoid noise.
- Use composite alarms to combine signals (error rate AND latency) to reduce false positives.
- Route alarm notifications through SNS to integrate with paging systems.

CDK snippet: alarm on Lambda error rate

```
import * as cw from 'aws-cdk-lib/aws-cloudwatch';

const metric = fn.metricErrors({ period: cdk.Duration.minutes(1) });
new cw.Alarm(this, 'LambdaErrorAlarm', { metric, threshold: 5,
evaluationPeriods: 3, alarmDescription: 'High lambda error rate' });
```

Security & regulatory

- Keep alarm history and attach runbooks. For financial services, store incidents and resolution steps for audits.

DR & Recovery

- Automate safe remediation (e.g., restart service) but gate production-critical automation with manual approval or progressive escalation.
-

8. Step Functions

Overview

Serverless orchestration. Standard (durable, audit-able) vs Express (high throughput, lower cost for short-lived flows).

Developer playbook & patterns

- Use Step Functions for transactional flows that need retries and clear state. Use explicit `Retry` and `Catch` clauses.
- Keep state small (avoid embedding large objects in state); use S3 pointers.
- Prefer integrated service actions (AWS SDK integration) to call AWS services directly from state machine (avoids Lambda shim).

CDK snippet: simple state machine

```
import * as sfn from 'aws-cdk-lib/aws-stepfunctions';
import * as tasks from 'aws-cdk-lib/aws-stepfunctions-tasks';

const hello = new tasks.LambdaInvoke(this, 'CallLambda', { lambdaFunction: fn,
payloadResponseOnly: true });
const sm = new sfn.StateMachine(this, 'StateMachine', { definition: hello,
stateMachineType: sfn.StateMachineType.STANDARD });
```

Observability

- Execution history is available in console; enable X-Ray for end-to-end tracing where needed.

Security & regulatory

- Control who can start/stop executions; log execution inputs/outputs carefully (avoid PII in plain text unless encrypted and audited).

DR & Recovery

- Replay capability for reprocessing fragments: if state machine supports idempotency and S3 pointers, you can re-run executions or run compensating transactions.

9. EventBridge

Overview

Central event bus for event-driven architectures. Supports schema registry, archive & replay, cross-account event delivery.

Developer playbook & patterns

- Define event schemas and version them; include `source`, `detailType`, `detail` JSON and correlation IDs.
- Use EventBridge archives to enable replay for recovery and testing.
- Use rule-level filtering to avoid fan-out storms.

CDK snippet: rule and target

```
import * as events from 'aws-cdk-lib/aws-events';
import * as targets from 'aws-cdk-lib/aws-events-targets';

const rule = new events.Rule(this, 'Rule', {
  eventPattern: { source: ['com.mycompany.orders'] }
});
rule.addTarget(new targets.LambdaFunction(fn));
```

Observability

- Monitor `Invocations`, `FailedInvocations`, `MatchedEvents` for rules, and check archive/replay counts.

Security & regulatory

- Use cross-account restrictions and resource policies to tightly control which principals can put events.
- Archive events to S3 for auditing and evidence in legal/regulatory events.

DR & Recovery

- Use EventBridge archive + replay to reprocess events into targets in another region/account after failover.

10. EC2

Overview

Traditional VM compute; used for workloads requiring full OS control, specialized drivers (GPUs), or custom networking.

Developer playbook & patterns

- Use AMI baking (Packer/Image Builder) to create golden images with security baseline.
- Use IMDSv2 enforced and instance roles for obtaining temporary credentials; do not store secrets on disk.
- Use Auto Scaling Groups with health checks and graceful termination hooks.

CDK snippet: simple auto-scaling group

```
import * as ec2 from 'aws-cdk-lib/aws-ec2';
import * as autoscaling from 'aws-cdk-lib/aws-autoscaling';

const vpc = ec2.Vpc.fromLookup(this, 'VPC', { isDefault: true });
const asg = new autoscaling.AutoScalingGroup(this, 'ASG', { vpc, instanceType:
  new ec2.InstanceType('t3.medium'), machineImage:
  ec2.MachineImage.latestAmazonLinux() });
```

Observability

- Install CloudWatch agent to gather system-level metrics (memory, disk) not available by default.

Security & regulatory

- Encryption of EBS, snapshots, strict IAM roles. Use VPC endpoints and private subnets for sensitive workloads.
- Patch management with SSM Patch Manager.

DR & Recovery

- Use AMI automation and ASG to recover compute quickly. For AZ/region level failures, deploy multi-region with DNS failover.

11. ECS (EC2 & Fargate)

Overview

Container orchestration — allow running containers on Fargate (serverless) or EC2 (self-managed host).

Developer playbook & patterns

- Use awsvpc network mode for task level networking and predictable security groups.
- For secrets, use Secrets Manager/SSM and reference in Task Definitions.

- Use health checks, deployment circuit breakers and ALB target group draining to ensure safe rollouts.

CDK snippet: Fargate service behind ALB

```
import * as ecs from 'aws-cdk-lib/aws-ecs';
import * as ecs_patterns from 'aws-cdk-lib/aws-ecs-patterns';

const cluster = new ecs.Cluster(this, 'Cluster', { vpc });
new ecs_patterns.ApplicationLoadBalancedFargateService(this, 'Service', {
  cluster, taskImageOptions: { image:
    ecs.ContainerImage.fromRegistry('myimage:latest') } });
```

Observability

- Collect container logs via awslogs or FireLens to central store; track task restarts and exit codes.

Security & regulatory

- Image scanning with ECR, use immutable tags, image signing where required.
- Isolate prod cluster in separate account or VPC for tenancy and compliance.

DR & Recovery

- Use multiple AZs for tasks. For region-level failover, deploy to alternate region and shift traffic via Route53.

Cross-cutting: Regulatory & Financial Services Considerations

Key controls to implement

- **Encryption:** all PII and financial data encrypted at rest (S3, EBS, DynamoDB, RDS) and in transit (TLS 1.2+). Use KMS CMKs when auditability & access control required. Rotate keys as per policy.
- **Auditability:** enable CloudTrail (management & data events), and log all API access. Keep immutable audit logs in a separate account and S3 bucket with object lock for legal hold.
- **Least privilege:** strict IAM policies, role separation (developer vs operator vs admin), and no long-lived credentials.
- **Data residency & sovereignty:** deploy to approved AWS regions; enforce via SCPs (AWS Organizations) and guardrails.
- **Retention & deletion:** implement data retention policies and secure deletion; document for compliance (SEC/FINRA etc.).

- **Change control:** use CI/CD pipelines with approvals for prod changes and maintain infrastructure-as-code history (Git).
- **Incident response / forensics:** centralize logs, enable packet captures where necessary, maintain runbooks and RACI.

Regulatory specifics to consider (asset management)

- Records retention (SEC/FINRA) — maintain transaction/audit records for mandated durations.
 - Trade surveillance — ensure event logging has required fields and timestamps.
 - Encryption & key management controls — proof of access controls and rotation.
 - Segregation of environments — production isolation and mandatory change approvals.
-

Cross-cutting: DR strategies & RTO/RPO recommendations

1. **Define RTO/RPO** per workload. Classify services: critical (RTO < 1 hour), important (RTO < 4 hours) and best-effort.
 2. **Multi-AZ first:** use AZ-resilient services by default (DynamoDB single-region multi-AZ, S3 multi-AZ durability, ECS multi-AZ tasks).
 3. **Multi-region** options:
 4. Use Global Tables (DynamoDB) and Cross-Region Replication (S3 CRR) for near-active-active or warm-standby.
 5. For compute/lambda, deploy to multiple regions and use Route53 weighted failover or Route53 health checks.
 6. **Event replay:** enable event archives (EventBridge) or store events in durable S3 + pointer patterns enabling replay.
 7. **Backups:** enable PITR (DynamoDB), automated snapshots (EBS), periodic AMI & RDS snapshots (if RDS used) and export logs to S3.
 8. **Runbooks & drills:** maintain runbooks and perform scheduled DR drills (failover + failback) with stakeholders.
-

Cross-cutting: Testing, CI/CD & Automation

- Use IaC (CDK, Terraform) for all infra — avoid manual console changes.
 - Build automated smoke tests and integration tests as part of pipelines; include replay tests for event-driven flows.
 - Use Canary deployments for production changes and automated rollbacks.
 - Add infrastructure tests (cfn-nag, cdk-nag, tfsec) to the pipeline.
 - Maintain separate AWS accounts for prod, staging, dev and centralize shared infra (logging/monitoring) in a separate account.
-

Example runbook template (use and adapt per integration)

1. **Title:** S3→SQS→Lambda processing failure
 2. **Symptoms:** Backlog in SQS, Lambda errors or DLQ fills.
 3. **Quick checks:**
 4. Check SQS ApproximateNumberOfMessagesVisible and AgeOfOldestMessage.
 5. Check Lambda CloudWatch logs for stack traces and throttles.
 6. Check DLQ for poison messages.
 7. **Remediation:**
 8. If code error: fix and redeploy via pipeline; reprocess DLQ items manually after verification.
 9. If throttling: scale consumers (Lambda concurrency, ECS task count) and/or increase visibility timeout.
 10. If messages duplicated: verify dedupe table TTL and error handling.
 11. **Postmortem:** record root cause, duration, and steps to prevent recurrence; update alarm thresholds if needed.
-

Attachments & next steps

- If you want, I can generate:
 - CDK repo skeleton (TypeScript) with setup pipeline and per-service examples.
 - Detailed runbook PDFs per integration.
 - IAM policy matrix for roles and responsibilities.
-

End of playbook.